

# Inverse Optimal Planning for Air Traffic Control

Ekaterina Tolstaya<sup>1†</sup>, Alejandro Ribeiro<sup>1</sup>, Vijay Kumar<sup>1</sup>, Ashish Kapoor<sup>2</sup>

<sup>1</sup> University of Pennsylvania

<sup>2</sup> Microsoft Corporation

† eig@seas.upenn.edu

International Conference on Intelligent Robots and Systems  
November 7, 2019

- ▶ Dense airspaces with:
  - ⇒ Human-controlled and autonomous aerial vehicles that co-exist safely
  - ⇒ Infrastructure for recreational, commercial and emergency aerial vehicles
- ▶ Autonomous air traffic control that imitates human air traffic control

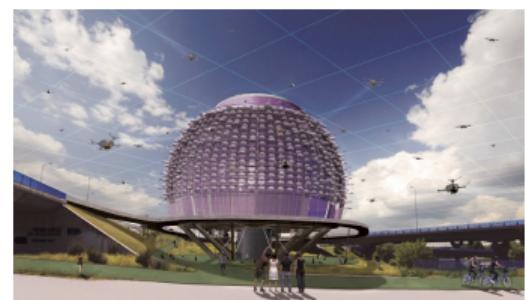


Reckless drone flies ABOVE passenger jet at McCarran Airport

(a) "Reckless drone flies ABOVE passenger jet" YouTube



(b) US Forest Service



(c) Architect Magazine

## Model-Based Planning

Dubins Airplane Model

Search-Based Planning

Dubins Airplane Heuristic

Encoding Safety and Preferences

## Learning from Demonstrations

Air Traffic Dataset

Maximum Entropy Inverse RL

Cost Function Parametrization

## Results

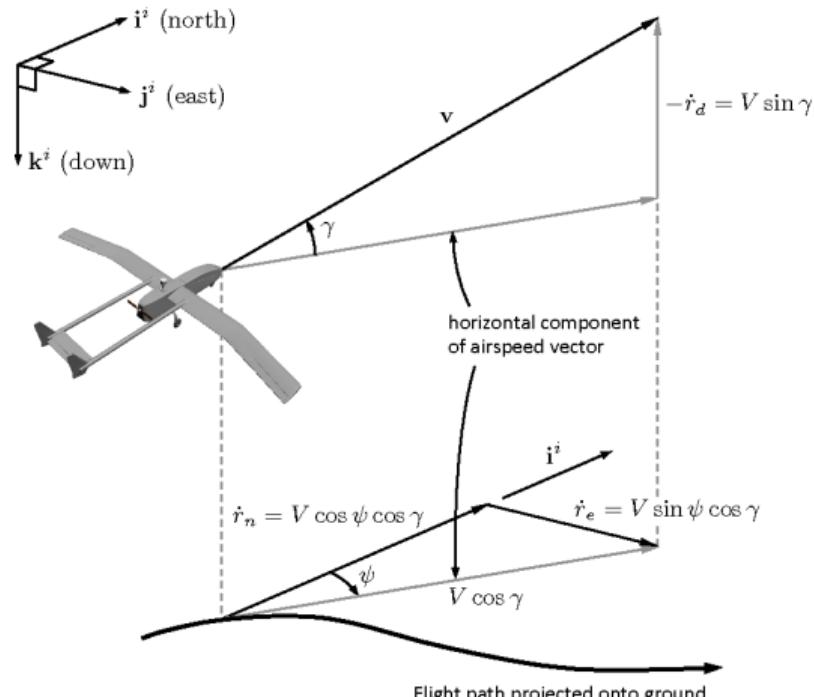
- ▶ Dubins airplane model,  $\mathbf{s} = (x, y, z, \phi)$

$$\dot{\mathbf{s}} = \begin{bmatrix} \dot{x} \\ \dot{y} \\ \dot{z} \\ \dot{\phi} \end{bmatrix} = \begin{bmatrix} v \cos \phi \\ v \sin \phi \\ u_z \\ u_\phi \end{bmatrix}$$

- ▶ Goal: plan a feasible trajectory

- ⇒ from an initial state  $\mathbf{s}_0 = (x_0, y_0, z_0, \phi_0)$
- ⇒ to a goal state  $\mathbf{s}_g = (x_g, y_g, z_g, \phi_g)$ .
- ⇒ that minimizes path length

$$\underset{\mathbf{s}(t)}{\operatorname{argmin}} \int_{t_0}^T \|\dot{\mathbf{s}}(t)\| dt$$

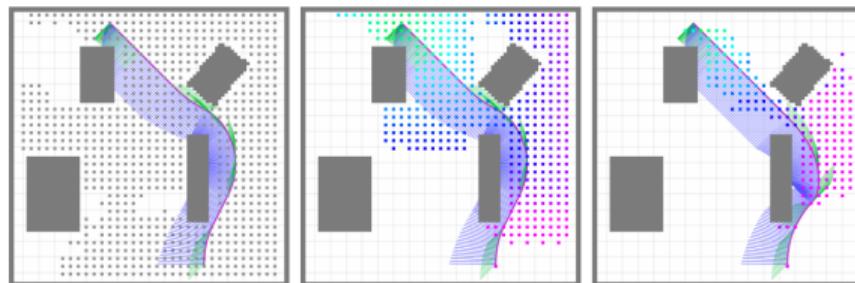


[OBM15]

- ▶ Search-based planning (versus sampling-based planning)
  - ⇒ Result is resolution-optimal (discretize space and time)
  - ⇒ Ability to deal with dynamic obstacles
  - ⇒ Ability to deal with non-holonomic systems and higher-order dynamics [LAMK17]
  - ⇒ Finite-time sub-optimality using the ARA\* search [LGT04]
- ▶ States are discretized with a resolution of  $\rho = [\rho_x, \rho_y, \rho_z, \rho_\phi]$ :

$$\bar{s} = \left[ \left\lfloor \frac{s_x}{\rho_x} \right\rfloor, \left\lfloor \frac{s_y}{\rho_y} \right\rfloor, \left\lfloor \frac{s_z}{\rho_z} \right\rfloor, \left\lfloor \frac{s_\phi}{\rho_\phi} \right\rfloor \right]$$

- ▶ Goal: find a path from  $s_0$  to  $s_g$  in graph  $\mathcal{G}$
- ▶ First try: Breadth-first search over all nodes in  $\mathcal{G}$  (Dijkstra's algorithm)
- ▶ Better: Prioritize nodes using a heuristic (A\*)
- ▶ Best: Anytime heuristic search (ARA\*) [LGT04]
  - ⇒ Over-inflates the heuristic to provide a sub-optimal solution in finite time
  - ⇒ Continually improves trajectory and reduces the sub-optimality bound



(a) Dijkstra.  $T_p = 0.16s, N_p = 2707$     (b)  $A^*$  with  $h_1$ .  $T_p = 0.064s, N_p = 1282$     (c)  $A^*$  with  $h_2$ .  $T_p = 0.016s, N_p = 376$

Figure. A\* and ARA\* require an admissible heuristic! [LAMK17]

**Algorithm 1** Dubins Airplane Heuristic

**Input:** Start  $s_0$ , goal  $s_g$ , forward speed  $v$ , max rate of climb  $\Delta z$ , turning rate  $\Delta\phi$

**Output:** Minimum path length from start to goal,  $d_{min}$

- 1: Compute Dubins car distance  $d_{xy}$  using the LSL, RSR, LSR, RSL, RLR, LRL paths with curvature  $\kappa = \Delta\phi/v$
- 2: Compute the minimum time from start to goal in the  $xy$  plane:  $t_{min} = d_{xy}/v$
- 3: Compute the minimum time for ascent/descent:
- 4:  $t_z = |z_g - z_0|/\Delta z$
- 5: **while**  $t_z > t_{min}$  **do**
- 6:     Add a helical ascent/descent to the trajectory:
- 7:      $t_{min} = t_{min} + \frac{2\pi}{\Delta\phi}$
- 8: **end while**
- 9: Compute the minimum path length from start to goal:  

$$d_{min} = \sqrt{(v \cdot t_{min})^2 + (z_g - z_0)^2}$$

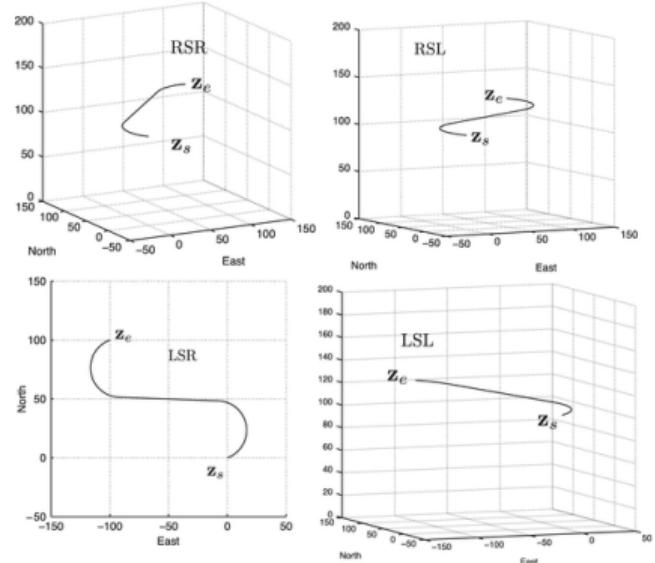


Figure. [OBM15]

# How to encode safety and requirements?

- ▶ Safety requirements result in a constrained optimization problem:

$$\begin{aligned} & \underset{\mathbf{s}(t)}{\operatorname{argmin}} \int_{t_0}^T \|\dot{\mathbf{s}}(t)\| dt \\ \text{s.t. } & \dot{\mathbf{s}}(t) = (\cos \phi, \sin \phi, u_z(t), u_\phi(t)) \\ & \mathbf{s}(t_0) = \mathbf{s}_0, \mathbf{s}(T) = \mathbf{s}_g \\ & \mathbf{s}(t) \in \mathcal{S}^{safe}, \mathbf{u}(t) \in \mathcal{U} \end{aligned} \tag{1}$$

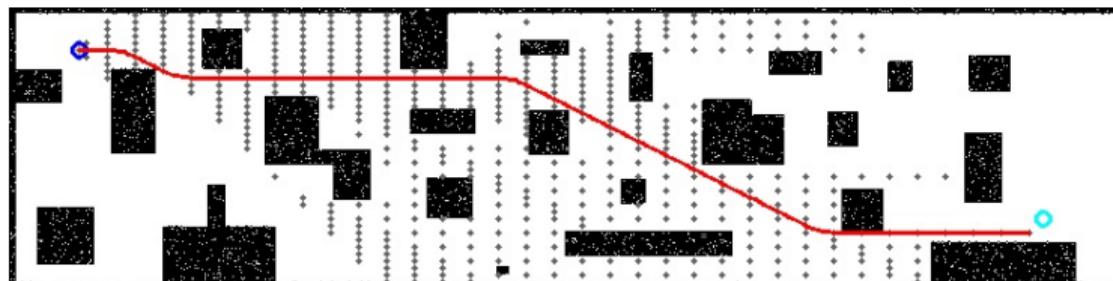


Figure. [LMAK18]

# How to encode safety and requirements?

- ▶ Potential function  $J$  to encode requirements and preferences of air traffic control
- ▶ The optimization problem is formulated as:

$$\operatorname{argmin}_{\mathbf{s}(t)} \int_{t_0}^T \left( 1 + J(\mathbf{s}(t)) \right) \|\dot{\mathbf{s}}(t)\| dt \quad (2)$$

$$\text{s.t. } \dot{\mathbf{s}}(t) = (\cos \phi, \sin \phi, u_z(t), u_\phi(t))$$

$$\mathbf{s}(t_0) = \mathbf{s}_0, \mathbf{s}(T) = \mathbf{s}_g$$

$$\mathbf{s}(t) \in \mathcal{S}^{\text{safe}}, \mathbf{u}(t) \in \mathcal{U}$$

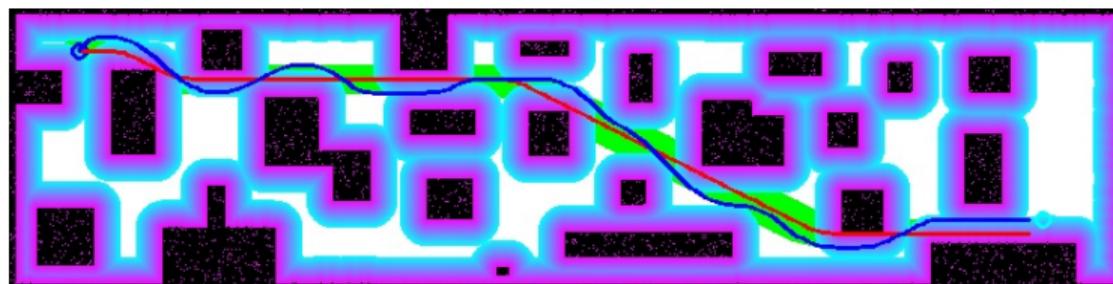
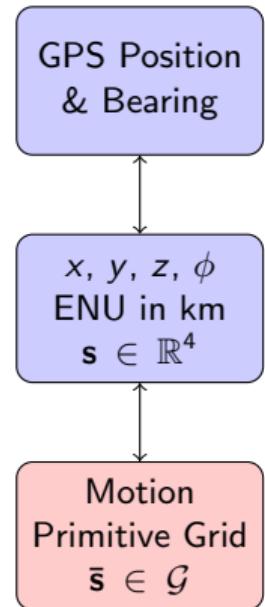
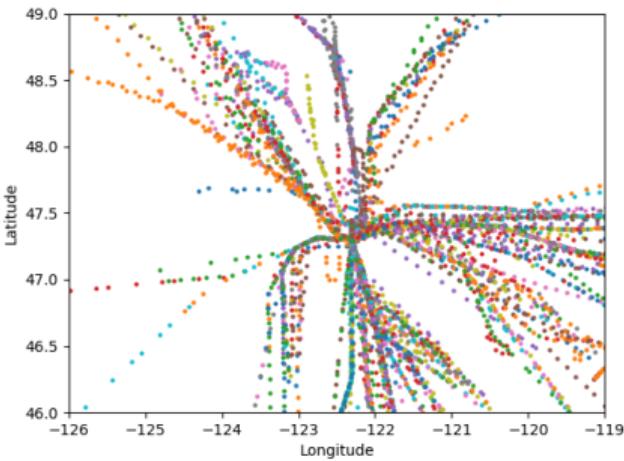
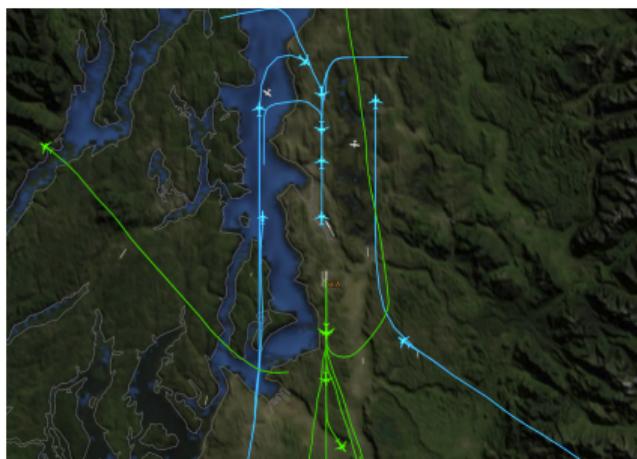


Figure. [LMAK18]

- ▶ Seattle-Tacoma airport data provided by FlightAware [Fli19]
  - ⇒ Local geography, air traffic density requirements
- ▶ Three parametrizations are used for expressing airplane states:
  - ⇒ Continuous state in blue, discrete - in red



- ▶ We are given a data set of **expert** demonstrations,  $\mathcal{D} = \{s_i^e(t)\}_{i=1,\dots,M}$
- ▶ Assumption: The expert demonstrations follow the principle of maximum entropy: [Jay57]:
  - ⇒ An expert is exponentially more likely to select a trajectory  $\tau$  with a lower cost:

$$\mathbb{P}(\tau) \propto e^{-C(\tau)} \tag{3}$$

- ▶ We are given a data set of **expert** demonstrations,  $\mathcal{D} = \{s_i^e(t)\}_{i=1,\dots,M}$
- ▶ Assumption: The expert demonstrations follow the principle of maximum entropy: [Jay57]:
  - ⇒ An expert is exponentially more likely to select a trajectory  $\tau$  with a lower cost:

$$\mathbb{P}(\tau) \propto e^{-C(\tau)} \quad (3)$$

- ▶ Hypothesis about the parametrization of the potential function

$$J_\theta(\mathbf{s}) = \theta^T \mathbf{f}(\mathbf{s}). \quad (4)$$

- ▶ Now, we seek to maximize the log likelihood of the expert trajectories:

$$\mathcal{L}(\theta) = \log \mathbb{P}(\mathcal{D}, \theta | J_\theta) \quad (5)$$

- ▶ Following the approach of [ZMBD08, WOP15]:

⇒ Expectation of the **learner's** state visitation counts:

$$\mathbb{E}[\mathbf{f}_\theta] = \sum_{\mathbf{s} \in \mathcal{S}} \mathbf{f}(\mathbf{s}) \mathbb{P}(\mathbf{s} | J_\theta) \quad (6)$$

⇒ **Expert's** empirical feature counts:

$$\mathbf{f}_{\mathcal{D}} = \sum_{\tau \in \mathcal{D}} \sum_{\mathbf{s} \in \tau} \mathbb{P}(\tau) \mathbf{f}(\mathbf{s}) \quad (7)$$

- ▶  $\nabla_\theta \mathcal{L}$  is the difference in feature counts of the trajectories of the **learner** and the **expert**

$$\nabla_\theta \mathcal{L} = \mathbb{E}[\mathbf{f}_\theta] - \mathbf{f}_{\mathcal{D}} \quad (8)$$

- ▶ Instead, we use the stochastic gradient computed using trajectories from the learner: [KPRS13]:

$$\hat{\nabla}_\theta \mathcal{L} = \underbrace{\frac{1}{N} \sum_{i=1}^N \sum_{t=t_0}^{T_i} \mathbf{f}(\mathbf{s}_i(t))}_{\text{learner}} - \underbrace{\frac{1}{M} \sum_{i=1}^M \sum_{t=t_0}^{T_i} \mathbf{f}(\mathbf{s}_i^e(t))}_{\text{expert}} \quad (9)$$

---

**Algorithm 2** Inverse Optimal Control for Air Traffic

---

- 1: **loop** until convergence
- 2:   Receive expert trajectories ordered by arrival time  $\mathcal{D} = \{\mathbf{s}_i^e(t)\}_{i=1,\dots,M}$
- 3:   **for**  $i = 0, 1, \dots, M$  **do**
- 4:     Obtain start state,  $\mathbf{s}_0$ , and end state,  $\mathbf{s}_g$ , of  $\mathbf{s}_i^e(t)$
- 5:     Obtain other airplanes' trajectories,  $\mathbf{s}_k^o(t)$  for  $k < i$
- 6:     Use ARA\* planner to minimize  $J_\theta$  over  $\mathbf{s}(t)$
- 7:     Compute stochastic gradient:

$$\hat{\nabla}_\theta \mathcal{L} = \sum_{t=t_0}^T f(\mathbf{s}(t)) - \sum_{t=t_0}^{T_i} f(\mathbf{s}_i^e(t))$$

- 8:       Update cost parameters  $\theta$  with step size  $\alpha$ :

$$\theta_{t+1} = \theta_t + \alpha \hat{\nabla}_\theta \mathcal{L}$$

- 9:       **end for**
  - 10:      **end loop**
-

- ▶ The motion cost must be:
  - ⇒ Non-negative
  - ⇒ Piece-wise linear in features of state  $f(s)$ , to satisfy the assumptions of [ZMBD08]

- ▶ The motion cost must be:
  - ⇒ Non-negative
  - ⇒ Piece-wise linear in features of state  $\mathbf{f}(\mathbf{s})$ , to satisfy the assumptions of [ZMBD08]
- ▶ The spatial potential field  $\bar{J}_a$  is approximated for each state  $\bar{\mathbf{s}} \in \mathcal{G}$ 
  - ⇒ Using a fine discretization of the input space

$$\bar{J}_a(\bar{\mathbf{s}}(t)) = \sum_{\bar{\mathbf{s}}_i \in \mathcal{G}} \max \{ \mathbf{w}_i, 0 \} \mathbb{1}_{\bar{\mathbf{s}}(t) = \bar{\mathbf{s}}_i}(\bar{\mathbf{s}}(t)), \quad (10)$$

- ▶ The motion cost must be:
  - ⇒ Non-negative
  - ⇒ Piece-wise linear in features of state  $\mathbf{f}(\mathbf{s})$ , to satisfy the assumptions of [ZMBD08]
- ▶ The spatial potential field  $\bar{J}_a$  is approximated for each state  $\bar{\mathbf{s}} \in \mathcal{G}$ 
  - ⇒ Using a fine discretization of the input space

$$\bar{J}_a(\bar{\mathbf{s}}(t)) = \sum_{\bar{\mathbf{s}}_i \in \mathcal{G}} \max \{ \mathbf{w}_i, 0 \} \mathbb{1}_{\bar{\mathbf{s}}(t) = \bar{\mathbf{s}}_i}(\bar{\mathbf{s}}(t)), \quad (10)$$

- ▶ Costs  $\mathbf{w}_i$  for each discrete state in the space are stored in a look-up table
  - ⇒ Enables more efficient motion planning

- ▶ To control the airspace around each airplane, we penalize pairwise distances between airplanes.
- ▶ Each prior arrival is treated as a moving cylinder with a known trajectory,  $\bar{s}_k^o(t)$

- ▶ To control the airspace around each airplane, we penalize pairwise distances between airplanes.
- ▶ Each prior arrival is treated as a moving cylinder with a known trajectory,  $\bar{s}_k^o(t)$
- ▶ We use a linear drop-off potential function [MMA00]:

$$\bar{J}_o(\bar{s}(t), \bar{s}_k^o(t)) = \mu \max\{\lambda_z - |\Delta_z(t)|, 0\} \cdot \max\{\lambda_{xy} - \sqrt{\Delta_x^2(t) + \Delta_y^2(t)}, 0\}, \quad (11)$$

- ▶  $\Delta(t) = \bar{s}(t) - \bar{s}_k^o(t)$  is the distance between two airplanes at time  $t$   
⇒  $\Delta_x(t)$ ,  $\Delta_y(t)$  and  $\Delta_z(t)$  are its components

- ▶ To control the airspace around each airplane, we penalize pairwise distances between airplanes.
- ▶ Each prior arrival is treated as a moving cylinder with a known trajectory,  $\bar{s}_k^o(t)$
- ▶ We use a linear drop-off potential function [MMA00]:

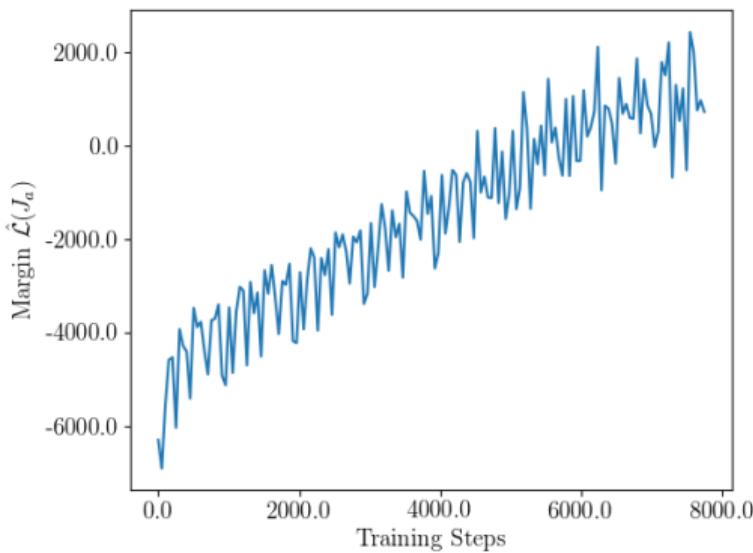
$$\bar{J}_o(\bar{s}(t), \bar{s}_k^o(t)) = \mu \max \{\lambda_z - |\Delta_z(t)|, 0\} \cdot \max \{\lambda_{xy} - \sqrt{\Delta_x^2(t) + \Delta_y^2(t)}, 0\}, \quad (11)$$

- ▶  $\Delta(t) = \bar{s}(t) - \bar{s}_k^o(t)$  is the distance between two airplanes at time  $t$   
 ⇒  $\Delta_x(t)$ ,  $\Delta_y(t)$  and  $\Delta_z(t)$  are its components
- ▶ Finally, the parameters for  $\bar{J}_a$  and  $\bar{J}_o$  can now be learned through stochastic gradient descent:

$$\theta = [\{w_i\}_{\bar{s}_i \in \mathcal{G}}, \lambda_{xy}, \lambda_z] \quad (12)$$

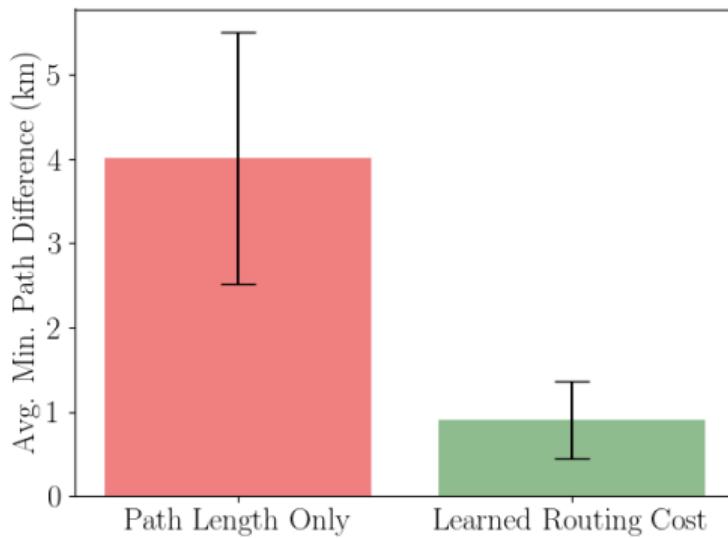
- ▶ The margin between the expert's cost and the learner's cost increases during training

$$\hat{\mathcal{L}}(\bar{J}_a) = \sum_{t=t_0}^{T_i} \bar{J}_a(\bar{s}(t)) - \sum_{t=t_0}^{T_i} \bar{J}_a(\bar{s}_i^e(t)) \quad (13)$$

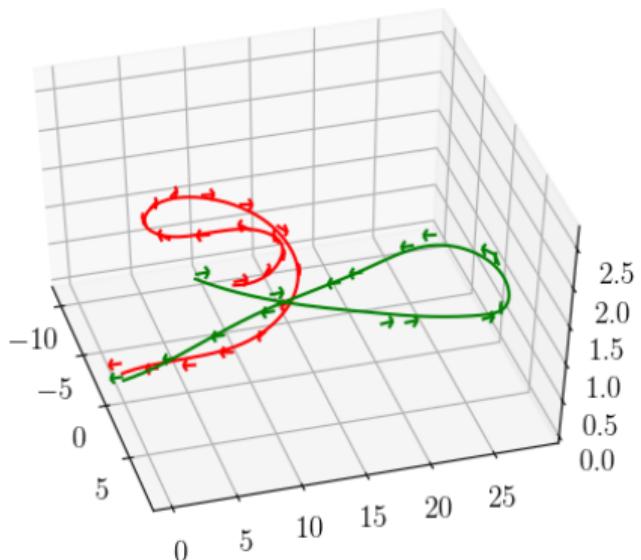


- ▶ The average minimum path difference
  - ⇒ How far the learner's trajectory strays from the expert's demonstration on average.

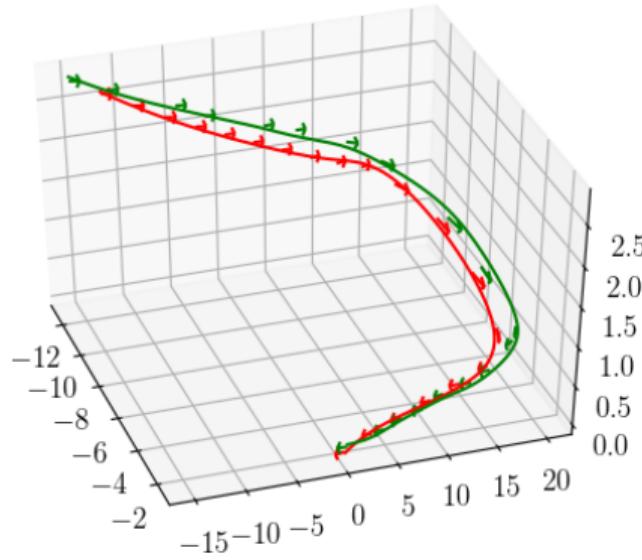
$$\frac{1}{N} \sum_{i=1}^N \frac{1}{T} \sum_{t=1}^T \min_u \|\bar{s}_i(t) - \bar{s}_i^e(u)\|_2$$



- ▶ A comparison of the expert's and learner's 3D trajectories for two different objective functions.
  - ⇒ The actual ATC trajectory is denoted with green arrows and a green spline.
  - ⇒ The planner's path is marked with red arrows and interpolated by the red curve.

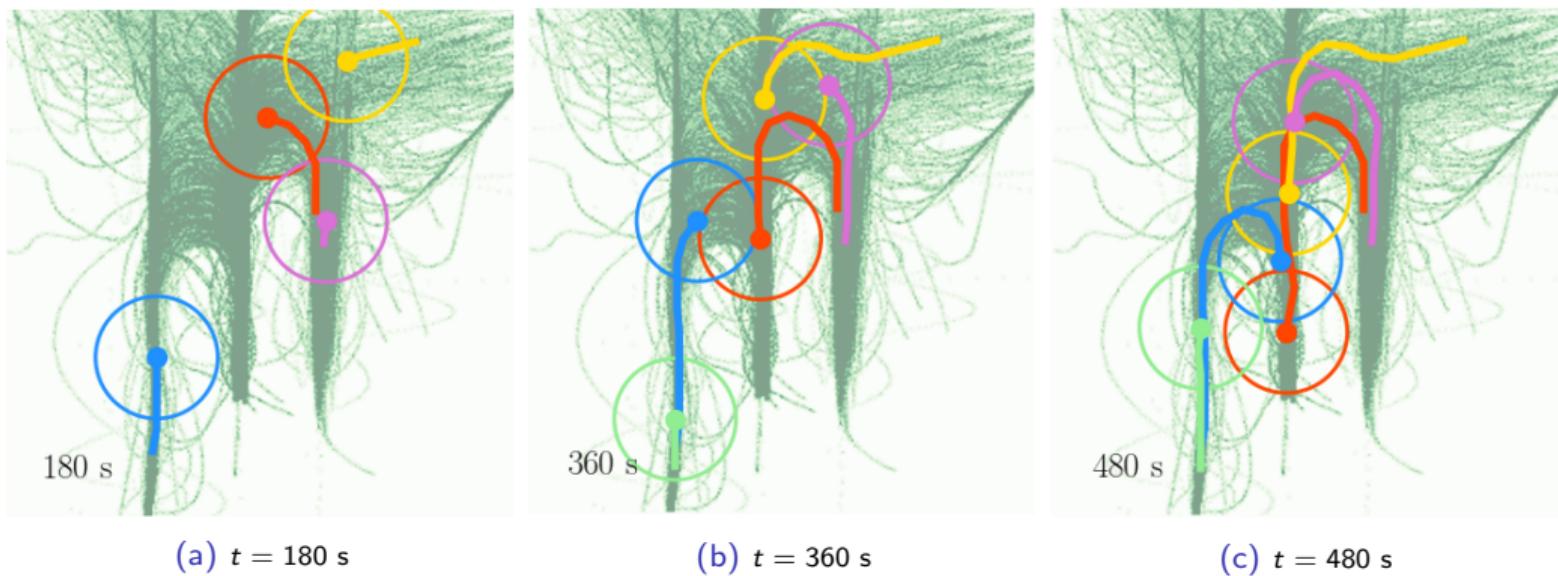


(a) Minimum path length

(b) Using  $\bar{J}_a$

# A multi-airplane scenario

- ▶ A top down view of five trajectories produced by the learner using both  $\bar{J}_a$  and  $\bar{J}_o$ .
- ▶ The large unfilled circles denote the threshold of the  $\bar{J}_o$  cost, maintaining separation
- ▶ The learned cost function  $\bar{J}_a$  is in the background indicated by green shading, with dark green indicating low cost states, and white indicating high cost states.



# Inverse Optimal Planning for Air Traffic Control

Ekaterina Tolstaya<sup>1†</sup>, Alejandro Ribeiro<sup>1</sup>, Vijay Kumar<sup>1</sup>, Ashish Kapoor<sup>2</sup>

<sup>1</sup> University of Pennsylvania

<sup>2</sup> Microsoft Corporation

† eig@seas.upenn.edu

International Conference on Intelligent Robots and Systems  
November 7, 2019

-  FlightAware, *Seattle-tacoma intl airport*, 2019.
-  Edwin T Jaynes, *Information theory and statistical mechanics*, Physical review **106** (1957), no. 4, 620.
-  Mrinal Kalakrishnan, Peter Pastor, Ludovic Righetti, and Stefan Schaal, *Learning objective functions for manipulation*, 2013 IEEE Int. Conf. on Robotics and Automation, IEEE, 2013, pp. 1331–1336.
-  Sikang Liu, Nikolay Atanasov, Kartik Mohta, and Vijay Kumar, *Search-based motion planning for quadrotors using linear quadratic minimum time control*, Intelligent Robots and Systems (IROS), 2017 IEEE/RSJ Int. Conf. on, IEEE, 2017, pp. 2872–2879.
-  Maxim Likhachev, Geoffrey J Gordon, and Sebastian Thrun, *Ara\*: Anytime a\* with provable bounds on sub-optimality*, Advances in neural information processing systems, 2004, pp. 767–774.
-  Sikang Liu, Kartik Mohta, Nikolay Atanasov, and Vijay Kumar, *Towards search-based motion planning for micro aerial vehicles*, arXiv preprint arXiv:1810.03071 (2018).
-  Robin Murphy, Robin R Murphy, and Ronald C Arkin, *Introduction to ai robotics*, MIT press, 2000.
-  Mark Owen, Randal W Beard, and Timothy W McLain, *Implementing dubins airplane paths on*